

# Weight Discretization for Quotient Model Abstraction in Spiking Neural Network Verification

*Research Note — January 2026*

## 1. Introduction

This note extends the filtration-based quotient model abstraction [1] to preserve *synaptic weight information* during formal verification. While the original quotient model (see `research_note_filtration.typ`) abstracts membrane potentials into equivalence classes, it treats all synapses uniformly, losing essential structural information.

We address this by introducing a *weight discretization scheme* that:

1. Maps continuous weights to a finite discrete range
2. Preserves the relative contribution of synapses to membrane potential
3. Ensures threshold feasibility is maintained
4. Retains weight visibility in the generated PRISM model

## 2. Weight Discretization

### 2.1. Formal Definition

**Definition.** Given a weight range  $[w_{\min}, w_{\max}] = [-100, 100]$  and a discretization parameter  $W \in \mathbb{N}^+$ , the *weight discretization function*  $\delta_W : \mathbb{Z} \rightarrow \mathbb{Z}$  is defined as:

$$\delta_{W(w)} = \text{round}\left(w \cdot \frac{W}{w_{\max}}\right)$$

The discretized weight range is  $[-W, W] \subset \mathbb{Z}$ .

**Example.** For  $W = 3$  (7 levels total):

- $\delta_3(100) = 3$  (strong excitatory)
- $\delta_3(67) = 2$  (medium excitatory)
- $\delta_3(33) = 1$  (weak excitatory)
- $\delta_3(0) = 0$  (negligible)
- $\delta_3(-50) = -2$  (medium inhibitory)
- $\delta_3(-100) = -3$  (strong inhibitory)

### 2.2. Threshold Calibration

The key challenge is ensuring that threshold reachability is preserved after discretization. We must calibrate the *discretized threshold*  $T_d$  to be consistent with the original threshold  $T$ .

**Definition.** The *discretized threshold* for a neuron with original threshold  $T$  and weight discretization parameter  $W$  is:

$$T_d = \text{ceil}\left(T \cdot \frac{W}{w_{\max}}\right)$$

**Theorem** (Threshold Preservation). Let  $\mathcal{N}$  be a neuron with incoming weights  $\{w_1, \dots, w_m\}$  and threshold  $T$ . Let  $\mathcal{N}'$  be the discretized version with weights  $\{\delta_W(w_1), \dots, \delta_W(w_m)\}$  and threshold  $T_d$ .

If  $\mathcal{N}$  can fire in a single step (i.e.,  $\exists$  input pattern  $\mathbf{y} \in \{0, 1\}^m$  such that  $\sum_{i=1}^m w_i \cdot y_i \geq T$ ), then  $\mathcal{N}'$  can also fire in a single step.

*Proof.* Let  $\mathbf{y}^*$  be an input pattern that causes  $\mathcal{N}$  to fire:

$$S = \sum_{i=1}^m w_i \cdot y_i^* \geq T$$

The discretized contribution is:

$$S_d = \sum_{i=1}^m \delta_W(w_i) \cdot y_i^*$$

By the rounding property of  $\delta_W$ :

$$\delta_W(w_i) \geq \frac{w_i \cdot W}{w_{\max}} - \frac{1}{2}$$

Therefore:

$$S_d \geq \left( \sum_{i=1}^m w_i \cdot y_i^* \right) \cdot \frac{W}{w_{\max}} - \frac{m}{2} = S \cdot \frac{W}{w_{\max}} - \frac{m}{2}$$

Since  $S \geq T$ :

$$S_d \geq T \cdot \frac{W}{w_{\max}} - \frac{m}{2}$$

For the discretized model to fire, we need  $S_d \geq T_d$ . By choosing:

$$T_d = \text{ceil}\left(T \cdot \frac{W}{w_{\max}}\right) \leq T \cdot \frac{W}{w_{\max}} + 1$$

A sufficient condition is:

$$S \cdot \frac{W}{w_{\max}} - \frac{m}{2} \geq T \cdot \frac{W}{w_{\max}} + 1$$

Which simplifies to:

$$(S - T) \cdot \frac{W}{w_{\max}} \geq \frac{m}{2} + 1$$

This holds when  $S - T$  is sufficiently large relative to  $m$ . For the boundary case  $S = T$ , we need  $W \geq w_{\max} \cdot (\frac{m}{2} + 1)/T$ .

In practice, with  $W = 3$ ,  $w_{\max} = 100$ , and typical networks ( $m \leq 10$ ,  $T = 100$ ):

$$W \geq 100 \cdot \frac{6}{100} = 6$$

Thus  $W = 7$  (range  $[-3, 3]$ ) is sufficient for most practical networks. For larger fan-in, a higher  $W$  may be required.  $\square$

**Remark.** The proof shows that weight discretization can introduce a *margin of error* proportional to the fan-in  $m$ . For high-fanin neurons ( $m > 10$ ), we recommend either:

1. Using a finer discretization ( $W \geq 5$ )
2. Applying a threshold correction factor:  $T_{d'} = T_d - \text{floor}\left(\frac{m}{2W}\right)$

### 3. Class Transition with Weighted Contributions

#### 3.1. Contribution-Based Class Evolution

In the original quotient model, class evolution used a binary rule:

- Any input fires  $\rightarrow$  class increases by 1
- No input fires  $\rightarrow$  class decreases by 1 (leak)

This loses weight information. We replace it with *weighted contribution-based* class evolution.

**Definition.** The *weighted contribution* for neuron  $n$  with incoming discretized weights  $\{w_1^d, \dots, w_m^d\}$  is:

$$C_n = \sum_{i=1}^m w_i^d \cdot y_i$$

where  $y_i \in \{0, 1\}$  is the spike output of presynaptic neuron  $i$ .

**Definition.** The *class delta function*  $\Delta : \mathbb{Z} \rightarrow \mathbb{Z}$  maps contribution to class change:

$$\Delta(C) = \text{clamp}\left(\text{round}\left(\frac{C}{\gamma}\right), -k, k\right)$$

where  $\gamma$  is the *class width* (typically  $\gamma = \frac{T_d}{k}$ ) and  $k$  is the number of threshold levels.

**Proposition.** The class delta function preserves the ordering of contributions: if  $C_1 > C_2$ , then  $\Delta(C_1) \geq \Delta(C_2)$ .

*Proof.* This follows directly from the monotonicity of `round` and `clamp` operations.  $\square$

#### 3.2. Integration with Leak Rate

The quotient model must also account for membrane potential decay (leak). We propose *discretized leak*:

**Definition.** The *discretized leak factor*  $\lambda_d$  is:

$$\lambda_d = 1 - \text{round}((1 - r) \cdot k)$$

where  $r \in [0, 1]$  is the original leak rate and  $k$  is the number of classes.

The class evolution rule becomes:

$$c'_n = \text{clamp}(c_n + \Delta(C_n) + \lambda_d, 0, k)$$

**Example.** For  $r = 0.9$  (10% decay per step) and  $k = 4$ :

$$\lambda_d = 1 - \text{round}(0.1 \cdot 4) = 1 - 0 = 1$$

So with no input, the class stays the same + leak = same - 0 (no change from this formula).

For  $r = 0.5$  (50% decay):

$$\lambda_d = 1 - \text{round}(0.5 \cdot 4) = 1 - 2 = -1$$

So with no input, class decreases by 1 per step.

**Remark.** We choose to apply leak *only when no excitatory input fires* to maintain consistency with the precise model's behavior, where leak is a multiplicative factor on the existing potential.

## 4. Threshold Feasibility Analysis

### 4.1. Definition

**Definition.** A neuron configuration is *threshold-feasible* if there exists at least one input pattern that can cause the neuron to reach the firing threshold within a finite number of steps.

**Theorem** (Feasibility Criterion). A neuron with discretized weights  $\{w_1^d, \dots, w_m^d\}$  and threshold  $T_d$  is threshold-feasible if and only if:

$$\sum_{w_i^d > 0} w_i^d \geq \frac{T_d}{1 + |\lambda_d|}$$

*Proof.* Let  $E = \sum_{w_i^d > 0} w_i^d$  be the maximum excitatory contribution per step.

*Sufficiency:* If  $E \geq T_d$ , the neuron can fire in a single step when all excitatory inputs fire. Otherwise, the neuron accumulates potential over steps. With leak factor  $\lambda_d \leq 0$ , each step adds at most  $E + \lambda_d$  net contribution.

For the class to reach  $k$  (firing threshold), starting from class 0:

$$n \cdot (E + \lambda_d) \geq k$$

The minimum steps required is  $n = \text{ceil}\left(\frac{k}{E + \lambda_d}\right)$ , which is finite iff  $E + \lambda_d > 0$ , i.e.,  $E > |\lambda_d|$ .

Since  $T_d = k \cdot \gamma$  and  $\gamma \approx \frac{T_d}{k}$ , the condition  $E \geq \frac{T_d}{1 + |\lambda_d|}$  ensures firing is reachable.

*Necessity:* If  $E < \frac{T_d}{1 + |\lambda_d|}$ , then even with optimal accumulation, the maximum class reached is bounded by  $\frac{E}{|\lambda_d|} < k$ , so firing is unreachable.  $\square$

## 4.2. Implementation

The feasibility check should be performed at PRISM generation time:

```
fn check_feasibility(
    weights: &[i32], // discretized weights
    threshold: i32,
    leak_factor: i32,
) -> Feasibility {
    let max_excitation: i32 = weights.iter()
        .filter(|&&w| w > 0)
        .sum();

    let min_required = threshold / (1 + leak_factor.abs());

    if max_excitation >= threshold {
        Feasibility::SingleStep
    } else if max_excitation >= min_required {
        let steps = (threshold + max_excitation - 1) / max_excitation;
        Feasibility::MultiStep { min_steps: steps }
    } else {
        Feasibility::Impossible
    }
}
```

## 5. PRISM Model Structure

The weighted quotient model generates PRISM code with explicit weight constants and contribution formulas:

```
// Weight constants (discretized)
const int W = 3; // Discretization parameter
const int W_in0_2 = 3; // δ_3(100) = 3
const int W_n1_2 = -2; // δ_3(-67) = -2
const int W_n3_2 = 1; // δ_3(33) = 1

// Discretized threshold
const int T_d = 3;

// Contribution formula (evaluated at runtime)
formula contrib_2 = W_in0_2 * x0 + W_n1_2 * z1_2 + W_n3_2 * z3_2;

// Class delta (clamped)
formula delta_2 = max(-4, min(4, contrib_2));

// Class evolution with weighted contribution
[tick] y2=0 & pClass2=0 -> (pClass2' = max(0, min(4, pClass2 + delta_2)));
```

## 6. Conclusion

We have established a formal framework for weight discretization in quotient model abstraction:

1. **Threshold Preservation Theorem** guarantees firing reachability is maintained
2. **Contribution-based class evolution** preserves weight effects on dynamics

3. **Feasibility analysis** detects and warns about unreachable configurations
4. **Discretized leak** maintains decay behavior in a state-preserving manner

The key insight is that weights enter the PRISM model as *constants*, not *state variables*, so weight discretization does not increase state space—it only affects transition probabilities and guard conditions.

---

## Bibliography

- [1] C. Baier and J.-P. Katoen, *Principles of Model Checking*. MIT Press, 2008.