

# Weight Discretization for Quotient Model Abstraction in Spiking Neural Network Verification

*Research Note — January 2026  
(Corrected & Revised)*

## 1. Introduction

This note extends the filtration-based quotient model abstraction [1] to preserve *synaptic weight information* during formal verification. While the original quotient model abstracts membrane potentials into equivalence classes, it treats all synapses uniformly, losing essential structural information.

We address this by introducing a *weight discretization scheme* that:

1. Maps continuous weights to a finite discrete range
2. Preserves the relative contribution of synapses to membrane potential
3. Ensures threshold feasibility is maintained
4. Retains weight visibility in the generated PRISM model

## 2. Weight Discretization

### 2.1. Formal Definition

**Definition.** Given a weight range  $[w_{\min}, w_{\max}] = [-100, 100]$  and a discretization parameter  $W \in \mathbb{N}^+$ , the *weight discretization function*  $\delta_W : \mathbb{Z} \rightarrow \mathbb{Z}$  is defined as:

$$\delta_{W(w)} = \text{round}\left(w \cdot \frac{W}{w_{\max}}\right)$$

The discretized weight range is  $[-W, W] \subset \mathbb{Z}$ .

**Example.** For  $W = 3$  (7 levels total):

- $\delta_3(100) = 3$  (strong excitatory)
- $\delta_3(67) = 2$  (medium excitatory)
- $\delta_3(33) = 1$  (weak excitatory)
- $\delta_3(0) = 0$  (negligible)
- $\delta_3(-50) = -2$  (medium inhibitory)
- $\delta_3(-100) = -3$  (strong inhibitory)

### 2.2. Threshold Calibration

The key challenge is ensuring that threshold reachability is preserved after discretization. We must calibrate the *discretized threshold*  $T_d$  to be consistent with the original threshold  $T$ .

**Definition.** The *discretized threshold* for a neuron with original threshold  $T$  and weight discretization parameter  $W$  is:

$$T_d = \text{ceil}\left(T \cdot \frac{W}{w_{\max}}\right)$$

**Theorem** (Threshold Preservation). Let  $\mathcal{N}$  be a neuron with incoming weights  $\{w_1, \dots, w_m\}$  and threshold  $T$ . Let  $\mathcal{N}'$  be the discretized version with weights  $\{\delta_W(w_1), \dots, \delta_W(w_m)\}$  and threshold  $T_d$ .

If  $\mathcal{N}$  can fire in a single step (i.e.,  $\exists$  input pattern  $\mathbf{y} \in \{0, 1\}^m$  such that  $\sum_{i=1}^m w_i \cdot y_i \geq T$ ), then  $\mathcal{N}'$  can also fire in a single step.

*Proof.* Let  $\mathbf{y}^*$  be an input pattern that causes  $\mathcal{N}$  to fire:

$$S = \sum_{i=1}^m w_i \cdot y_i^* \geq T$$

The discretized contribution is:

$$S_d = \sum_{i=1}^m \delta_W(w_i) \cdot y_i^*$$

By the rounding property of  $\delta_W$ :

$$\delta_W(w_i) \geq \frac{w_i \cdot W}{w_{\max}} - \frac{1}{2}$$

Therefore:

$$S_d \geq \left( \sum_{i=1}^m w_i \cdot y_i^* \right) \cdot \frac{W}{w_{\max}} - \frac{m}{2} = S \cdot \frac{W}{w_{\max}} - \frac{m}{2}$$

Since  $S \geq T$ :

$$S_d \geq T \cdot \frac{W}{w_{\max}} - \frac{m}{2}$$

For the discretized model to fire, we need  $S_d \geq T_d$ . By choosing:

$$T_d = \text{ceil}\left(T \cdot \frac{W}{w_{\max}}\right) \leq T \cdot \frac{W}{w_{\max}} + 1$$

A sufficient condition is:

$$S \cdot \frac{W}{w_{\max}} - \frac{m}{2} \geq T \cdot \frac{W}{w_{\max}} + 1$$

Which simplifies to:

$$(S - T) \cdot \frac{W}{w_{\max}} \geq \frac{m}{2} + 1$$

This holds when  $S - T$  is sufficiently large relative to  $m$ . For the boundary case  $S = T$ , we need  $W \geq w_{\max} \cdot (\frac{m}{2} + 1)/T$ .

In practice, with  $W = 3$ ,  $w_{\max} = 100$ , and typical networks ( $m \leq 10$ ,  $T = 100$ ):

$$W \geq 100 \cdot \frac{6}{100} = 6$$

Thus  $W = 7$  (range  $[-3, 3]$ ) is sufficient for most practical networks.  $\square$

**Remark. Critical Constraint:** The proof shows that weight discretization introduces a cumulative error of  $-\frac{m}{2}$ . If the fan-in  $m$  is large relative to  $W$  (specifically if  $m > 2W$ ), the rounding noise may exceed the signal of the smallest synaptic weight.

For high-fanin neurons, we strictly recommend:

1. Using a finer discretization ( $W \geq \frac{m}{2}$ )
2. Applying a threshold correction factor:  $T_{d'} = T_d - \text{floor}\left(\frac{m}{2W}\right)$  (Note: This may increase false positive firings).

### 3. Class Transition with Weighted Contributions

#### 3.1. Contribution-Based Class Evolution

In the original quotient model, class evolution used a binary rule which loses weight information. We replace it with *weighted contribution-based* class evolution.

**Definition.** The *weighted contribution* for neuron  $n$  with incoming discretized weights  $\{w_1^d, \dots, w_m^d\}$  is:

$$C_n = \sum_{i=1}^m w_i^d \cdot y_i$$

where  $y_i \in \{0, 1\}$  is the spike output of presynaptic neuron  $i$ .

**Definition.** The *class delta function*  $\Delta : \mathbb{Z} \rightarrow \mathbb{Z}$  maps contribution to class change:

$$\Delta(C) = \text{clamp}\left(\text{round}\left(\frac{C}{\gamma}\right), -k, k\right)$$

where  $\gamma$  is the *class width* (typically  $\gamma = \frac{T_d}{k}$ ) and  $k$  is the number of threshold levels.

#### 3.2. Integration with Leak Rate

The quotient model must also account for membrane potential decay (leak). We propose a corrected *discretized leak* formula.

**Definition.** The *discretized leak factor*  $\lambda_d$  is:

$$\lambda_d = -\text{round}((1 - r) \cdot k)$$

where  $r \in [0, 1]$  is the original leak rate (retention) and  $k$  is the number of classes.

The class evolution rule becomes:

$$c'_n = \text{clamp}(c_n + \Delta(C_n) + \lambda_d, 0, k)$$

**Example.** For  $r = 0.9$  (90% retention, i.e., 10% decay) and  $k = 4$ :

$$\lambda_d = -\text{round}((1 - 0.9) \cdot 4) = -\text{round}(0.4) = 0$$

*Result:* With no input, the class stays stable (decay is too small to drop a full class).

For  $r = 0.5$  (50% decay):

$$\lambda_d = -\text{round}((1 - 0.5) \cdot 4) = -\text{round}(2.0) = -2$$

*Result:* With no input, class decreases by 2 per step.

**Remark.** We choose to apply leak *only when no excitatory input fires* to maintain consistency with the precise model's behavior, where leak is a multiplicative factor on the existing potential.

## 4. Threshold Feasibility Analysis

### 4.1. Definition

**Definition.** A neuron configuration is *threshold-feasible* if there exists at least one input pattern that can cause the neuron to reach the firing threshold within a finite number of steps.

**Theorem** (Feasibility Criterion). A neuron with discretized weights  $\{w_1^d, \dots, w_m^d\}$  and threshold  $T_d$  is threshold-feasible if and only if:

$$\sum_{w_i^d > 0} w_i^d > |\lambda_d|$$

And specifically, for reliable firing:

$$\sum_{w_i^d > 0} w_i^d \geq \frac{T_d}{1 + |\lambda_d|}$$

*Proof.* Let  $E = \sum_{w_i^d > 0} w_i^d$  be the maximum excitatory contribution per step.

*Sufficiency:* If  $E \geq T_d$ , the neuron can fire in a single step. Otherwise, the neuron accumulates potential. With leak factor  $\lambda_d \leq 0$ , each step adds  $E + \lambda_d$  net contribution.

For accumulation to be possible at all, we must have  $E + \lambda_d > 0$ , or  $E > |\lambda_d|$ .

If this holds, the minimum steps required to reach class  $k$  is  $n = \text{ceil}\left(\frac{k}{E + \lambda_d}\right)$ . Since  $T_d = k \cdot \gamma$ , the condition  $E \geq \frac{T_d}{1 + |\lambda_d|}$  ensures firing is reachable.

*Necessity:* If  $E \leq |\lambda_d|$ , the potential cannot grow over time; the leak cancels or overpowers the excitation.  $\square$

### 4.2. Implementation

The feasibility check should be performed at PRISM generation time:

```

fn check_feasibility(
    weights: &[i32], // discretized weights
    threshold: i32,
    leak_factor: i32, // expected to be <= 0 (e.g. -1, -2)
) -> Feasibility {
    let max_excitation: i32 = weights.iter()
        .filter(|&&w| w > 0)
        .sum();

    // Ensure we are working with the magnitude of the leak
    let leak_magnitude = leak_factor.abs();

    // Basic sanity check: Input must overcome leak
    if max_excitation <= leak_magnitude {
        return Feasibility::Impossible;
    }

    let min_required = threshold / (1 + leak_magnitude);

    if max_excitation >= threshold {
        Feasibility::SingleStep
    } else if max_excitation >= min_required {
        // Steps = ceil(Threshold / Net_Gain)
        let net_gain = max_excitation - leak_magnitude;
        let steps = (threshold + net_gain - 1) / net_gain;
        Feasibility::MultiStep { min_steps: steps }
    } else {
        Feasibility::Impossible
    }
}

```

## 5. PRISM Model Structure

The weighted quotient model generates PRISM code with explicit weight constants and contribution formulas. Note that weights are static constants, minimizing state explosion.

```

// Weight constants (discretized)
const int W = 3; // Discretization parameter
const int W_in0_2 = 3; // δ_3(100) = 3
const int W_n1_2 = -2; // δ_3(-67) = -2
const int W_n3_2 = 1; // δ_3(33) = 1

// Discretized threshold
const int T_d = 3;

// Contribution formula (evaluated at runtime)
formula contrib_2 = W_in0_2 * x0 + W_n1_2 * z1_2 + W_n3_2 * z3_2;

// Class delta (clamped)
formula delta_2 = max(-4, min(4, contrib_2));

// Class evolution with weighted contribution
// Note: Leak is applied via addition of negative constant lambda_d
[tick] y2=0 & pClass2=0 -> (pClass2' = max(0, min(4, pClass2 + delta_2)));

```

## 6. Conclusion

We have established a formal framework for weight discretization in quotient model abstraction. The key corrections in this version include:

1. **Corrected Leak Formulation:** Using ensures leak acts as a decay, not a driver.
  2. **Fan-in Awareness:** Highlighting that must scale with fan-in to prevent rounding errors from dominating dynamics.
  3. **Feasibility Logic:** Ensuring excitation strictly exceeds leak magnitude () for multi-step firing.
- 

## Bibliography

- [1] C. Baier and J.-P. Katoen, *Principles of Model Checking*. MIT Press, 2008.